

// CONTACT

- 07734 258 202
- josepholman2k@tuta.com
- London, UK
- linkedin
- website

// AI & ML

LLMs

RAG

LangChain

LangGraph

MCP

Generative AI

Prompt Eng.

PyTorch

TensorFlow

MLflow

Ragas

Hugging Face

scikit-learn

LoRA

Embeddings

NLP

NumPy

// DATA & INFRA

Kafka

Elasticsearch

MongoDB

PostgreSQL

Redis

Pandas

Snowflake

Dramatiq

Kendra

Terraform

GitHub Actions

AWS Bedrock

// CLOUD & DEVOPS

AWS

Docker

Kubernetes

HelM

Jenkins

Grafana

Argo CD

CI/CD

// DEVELOPMENT

Python

FastAPI

Pydantic

JavaScript

REST APIs

SQL

asyncio

Linux

Bash

Pytest

Git

Agile

// OBSERVABILITY

Grafana

Prometheus

OpenTelemetry

Sentry

CloudWatch

Structlog

ELK

// EDUCATION

BSc Computer Science
& Artificial Intelligence
University of Sussex · 2023
OFirst Class Honours (1:1)

// WORK AUTH

 British Citizen

Eligible for EU Blue Card (IT shortage occupation). Available to relocate within the EU on standard 2-3 month Blue Card timeline.

\$ whoami

Joseph Holman

AI Engineer with a strong core in LLM applications, RAG systems and agentic architecture. I like to work on tough problems with like-minded people.

EXPERIENCE

TryHackMe

Apr 2025 - Present

AI Engineer, Remote

- As a founding member of the AI team I was responsible for researching, building and deploying AI features inline with the founders and product vision.
- Using user VM telemetry I shipped an adaptive AI tutor. Designed a LangGraph agent workflow that interprets raw Linux/ Windows session data, decides whether the learner needs a hint, correction or no help at all. Drove a 25% lift in question to completion vs unassisted exercises.
- Built an AI feedback system comparing user telemetry against benchmarks to surface knowledge gaps and generate personalised learning paths, users seeing this feature are 15% more likely to attempt a follow-up exercise increasing time spent on the platform.
- Built + deployed production AI services from prototypes to 200k+ MAU. Designed event-driven infra on AWS EKS (MongoDB, Redis, Snowflake, Elasticsearch).
- Built and shipped a content recommendation engine combining vector embedding, collaborative filtering and real-time user data into a configurable ranker. Backed with an offline+online evaluation pipeline tracking ranking metrics that gates releases and auto flags regressions before they ship.
- Designed evaluation pipeline (LLM-as-judge + Ragas) catching regressions before prod, enabling safe A/B shadow rollouts.
- Mentoring a software engineer transitioning into AI, scoping and delegating tasks while building their capability across the team's AI features.

bp

Dec 2023 - Jan 2025

Software Engineer, London (Contract)

- Embedded with bp as a contractor to deliver their first production generative RAG system. Ran discovery with documentation and product owners to scope what "good retrieval" meant against their internal corpus.
- Built custom Kendra data source connector to integrate internal repos, tuned relevance against user queries, exposed faithfulness, relevancy and context precision through MLflow so non-technical stakeholders could review readiness themselves.
- Extended into automated repository monitoring validating API and CICD docs against company standards after initial RAG solution surfaced the underlying problem.
- Implemented AWS synthetic monitoring (CloudWatch) for UI testing and proactive alerting.

FutureNow

Sep 2023 - Feb 2025

Software Engineer, London (Part-time)

- Implemented FastAPI backend services for an AI-driven knowledge management product, using LangChain to automate code documentation generation.
- Re-architected a single-agent LLM system into a multi-agent design using LangGraph with PostgreSQL-backed state, improving scalability and response accuracy.
- Designed evaluation pipelines monitoring real-time AI metrics.